

Planning in Factored Spaces

Tom Silver

Robot Planning Meets Machine Learning

Princeton University

Fall 2025

Recap and Preview

Previously:

- Planning in finite “tabular” state and action spaces
- Careful treatment of uncertainty in transitions and observations
- Offline planning and online planning

Now:

- Planning in finite “**factored**” state and action spaces
- **No more uncertainty**
- **Online planning** only




Our focus turns to leveraging
structure in the problem space

Later:

- Planning in **continuous** state and action spaces

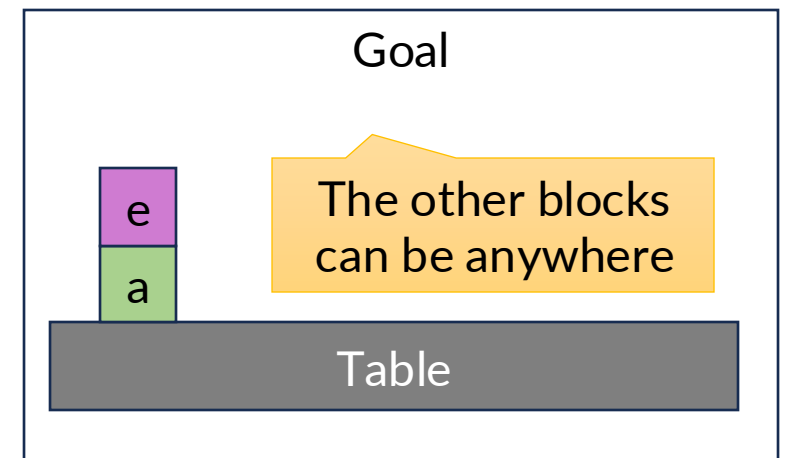

Classical Planning Problem Setting

A classical planning problem is:

1. A finite state space \mathcal{S}
2. A finite action space \mathcal{A}
3. An initiable action function $I: \mathcal{S} \times \mathcal{A} \rightarrow \{T, F\}$
4. A transition function $F: \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$  Deterministic! Can be partial
5. A cost function $C: \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}_{\geq 0}$  Lower better. Could do rewards instead; just a convention.
6. An initial state $s_0 \in \mathcal{S}$
7. A goal function $G: \mathcal{S} \rightarrow \{T, F\}$  Equivalent to a set of states

Example: Blocks World

- **States:** each block is either on the table or on some other block
- **Actions:** picking or placing a block
- **Initiation:** can only pick and place on “clear” blocks
- **Transition function:** as you’d expect
- **Cost function:** always 1
- **Initial state:** e.g., see right
- **Goal function:** e.g., see right



Definition of a Solution (Plan)

A *solution* ψ to a classical planning problem is a sequence of states s_0, s_1, \dots, s_T and actions a_0, a_1, \dots, a_{T-1} such that

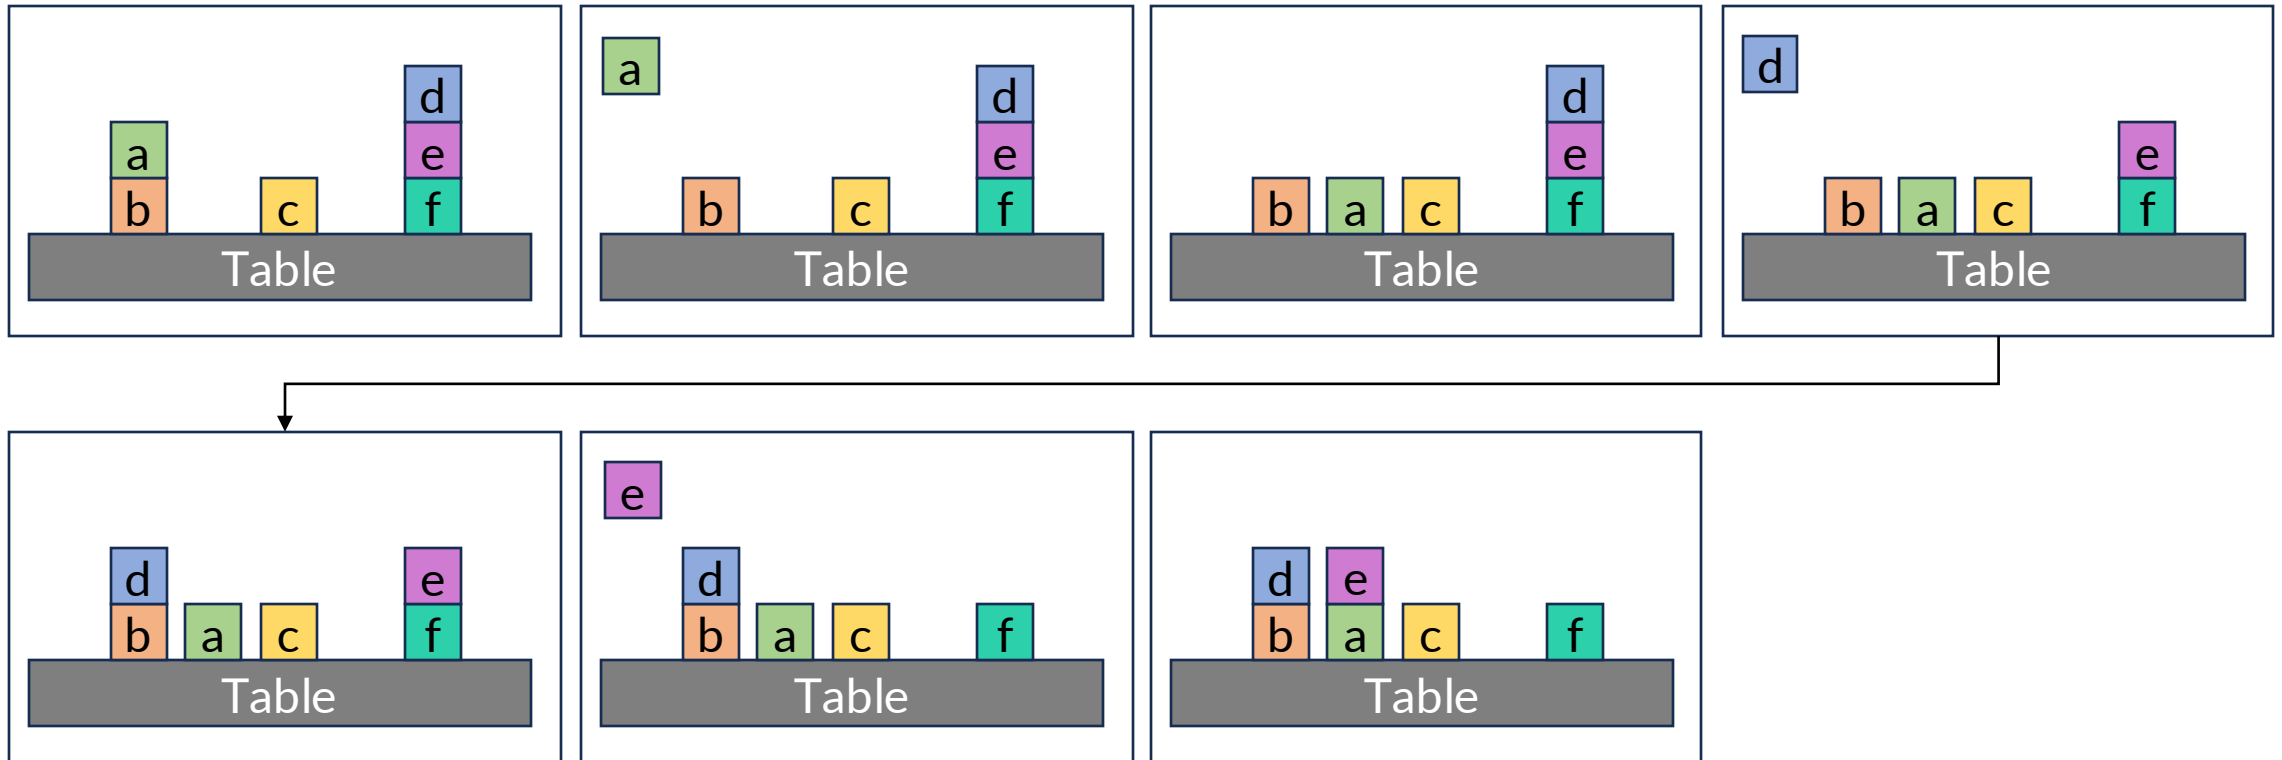

1. Each action is initiatable: $I(s_t, a_t) = \text{True}$
2. Transitions are valid: $T(s_t, a_t) = s_{t+1}$
3. The goal is achieved: $G(s_T) = \text{True}$

The *cost* of a solution ψ is $C(\psi) \triangleq \sum_t C(s_t, a_t, s_{t+1})$

A solution ψ^* is *optimal* if it minimizes costs: $C(\psi^*) = \min_{\psi} C(\psi)$

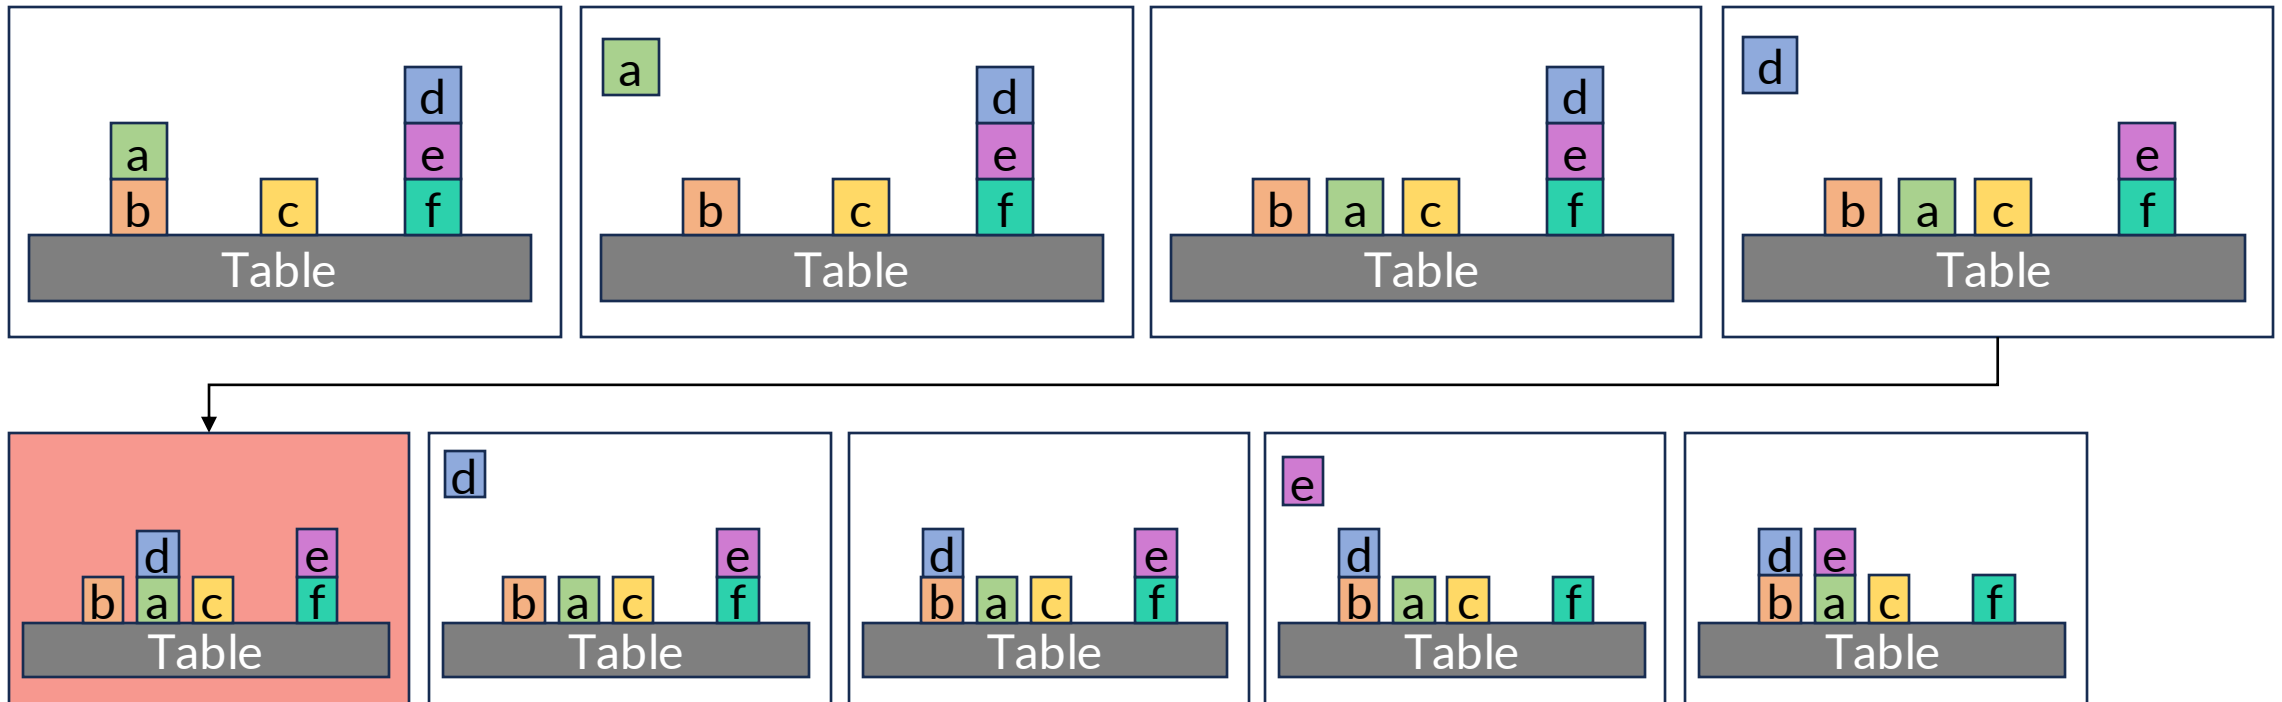

Example: Blocks World

An optimal plan



Example: Blocks World

A **sub**optimal plan



A Stupidest Possible Algorithm

Randomly sample applicable actions until the goal is reached.

A Stupidest Possible Algorithm

Randomly sample applicable actions until the goal is reached.

Definitions:


A planner is *sound* if its output is guaranteed to be a solution.

A planner is *complete* if it is guaranteed to return an output eventually.

A sound planner is *optimal* if its output is guaranteed to be optimal.
Otherwise, the planner is *satisficing*.

What is this planner?

A Better Approach: Graph Search




Graph Search

GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

Goal: b on a



GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0$ , parent=null)
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c'$ , parent=node)
20        push child onto queue with PRIORITY(child)
```


queue



GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = Node( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = Node( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue


(, 0)

For now, let's say
priority = path
cost


GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0$ , parent=null)
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c'$ , parent=node)
20        push child onto queue with PRIORITY(child)
```

queue

(, 0)

bestPathCost




→ 0

GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost



→ 0




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10      $s, c = \text{node.state}, \text{node.pathCost}$ 
11     // skip if we already found a better path
12     if bestPathCost[s] < c : continue
13     if  $G(s)$  : return node.extractPlan()
14     for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15          $s' = F(s, a)$ 
16          $c' = c + C(s, a, s')$ 
17         if bestPathCost[s'] <  $c'$  : continue
18         bestPathCost[s'] =  $c'$ 
19         child = NODE( $s', c', \text{parent}=\text{node}$ )
20         push child onto queue with PRIORITY(child)
```

queue

bestPathCost



→ 0




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost



→ 0




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost



→ 0




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)


```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```


queue

bestPathCost



→ 0




GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state → best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```


queue

bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

$(\text{○}, 1)$

bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

$(\text{orange circle}, 1)$

bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

$(\text{orange circle}, 1) (\text{blue circle}, 1)$


bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

(, 1)


bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] < c : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```


queue

(, 1)


bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

(, 1)

bestPathCost





Path cost is 2, which is worse than 0



GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```


queue

(, 1) (, 2)




bestPathCost





GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

queue

(, 1) (, 2) (, 2)

bestPathCost



GRAPHSEARCH($s_0, \mathcal{A}, I, F, C, G$)

```
1 // priority queue of nodes
2 initialize queue = []
3 root = NODE( $s_0, 0, \text{parent}=\text{null}$ )
4 // PRIORITY differs between algorithms
5 push root onto queue with PRIORITY(root)
6 // state  $\rightarrow$  best known path cost
7 initialize bestPathCost =  $\{s_0 \mapsto 0\}$ 
8 while queue is not empty
9     pop node from queue
10     $s, c = \text{node.state}, \text{node.pathCost}$ 
11    // skip if we already found a better path
12    if bestPathCost[s] <  $c$  : continue
13    if  $G(s)$  : return node.extractPlan()
14    for  $a \in \mathcal{A}$  s.t.  $I(s, a)$ 
15         $s' = F(s, a)$ 
16         $c' = c + C(s, a, s')$ 
17        if bestPathCost[s'] <  $c'$  : continue
18        bestPathCost[s'] =  $c'$ 
19        child = NODE( $s', c', \text{parent}=\text{node}$ )
20        push child onto queue with PRIORITY(child)
```

Heuristics and Value Functions

As in MDP land, we can define value functions:

$$V^*(s) = \begin{cases} 0 & \text{if } G(s) \\ \min_{\substack{a: I(s,a) \\ s' = F(s,a)}} C(s, a, s') + V^*(s') & \text{o. w.} \end{cases}$$

“Cost-to-go”

Heuristics and Value Functions

As in MDP land, we can define value functions:

$$V^*(s) = \begin{cases} 0 & \text{if } G(s) \\ \min_{\substack{a: I(s,a) \\ s'=F(s,a)}} C(s, a, s') + V^*(s') & \text{o. w.} \end{cases}$$

A **heuristic** $\hat{V}(s)$ is an approximate value function.

Same as MDP land

Heuristics and Value Functions

As in MDP land, we can define value functions:

$$V^*(s) = \begin{cases} 0 & \text{if } G(s) \\ \min_{\substack{a: I(s,a) \\ s'=F(s,a)}} C(s, a, s') + V^*(s') & \text{o. w.} \end{cases}$$

A **heuristic** $\hat{V}(s)$ is an approximate value function.

A heuristic is **admissible** if it never overestimates the cost-to-go:

$$\text{For all } s \in \mathcal{S}, \hat{V}(s) \leq V^*(s).$$

Graph Search Variations

Algorithm	Priority Function	Optimal?	Notes
Uniform cost search	path cost	Yes	If costs are 1, this is breadth-first search. Like Dijkstra's, but returns shortest path to goal, not shortest paths to all states
Greedy best-first search (GBFS)	heuristic (state)	No	Good choice for fast satisficing planning
A* search	path cost + heuristic (state)	Depends	Optimal if heuristic is <i>admissible</i> (never overestimates cost-to-go)
Depth first search	negative path cost	No	Can be more memory-efficient if implemented as a special case

Where do Heuristics Come From?

1. Hand-designed based on understanding of the problem
2. Learned from data (later in the course)
3. **Automatically derived from the problem representation**

Factored Classical Planning Problems

Consider a classical planning problem where:

States are *factored* into n Boolean features:


$$\mathcal{S} = \{T, F\}^n$$

The goal is to “activate” features $\{i_1, \dots, i_m\}$ (for $1 \leq i_j \leq n$):

$$G(s) = s[i_1] \wedge \dots \wedge s[i_m]$$

Blocks World Example

Feature	Value
On-A-B	True
On-A-C	False
On-B-A	False
On-B-C	False
On-C-A	False
On-C-B	False
OnTable-A	False
OnTable-B	True
OnTable-C	True
Holding-A	False
Holding-B	False
Holding-C	False
HandEmpty	True





Blocks World Example

Feature	Value
On-A-B	False
On-A-C	False
On-B-A	False
On-B-C	False
On-C-A	False
On-C-B	False
OnTable-A	True
OnTable-B	True
OnTable-C	True
Holding-A	False
Holding-B	False
Holding-C	False
HandEmpty	True

Goal: a on b &
b on c



{On-A-B,
On-B-C}





Goal-Count: Our First Problem-Derived Heuristic

The goal-count heuristic counts the number of goal features that are not yet activated:

$$V_{GC}(s) \triangleq |\{i : \neg s[i] \wedge i \in G\}|$$

Assuming all transition costs are 1,
is V_{GC} admissible?

Goal-Count Can Help!



Limitations of Goal-Count

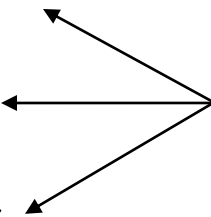
1. Very sparse

2. Can be “misleading”

Examples?

Factoring Further: Actions + Transitions

A (STRIPS / PDDL) *operator* has:

1. Preconditions
 2. Add effects
 3. Delete effects
- 
- Each is a set of features

Notation: ω is an operator, Ω is the set of all operators

Pick-A-from-C:

Preconditions: {HandEmpty,
On-A-C,
Clear-A}

Add effects: {Holding-A, Clear-C}

Delete effects: {HandEmpty,
On-A-C, Clear-A}

Factored Classical Planning Problems

A *factored classical planning problem* is:

1. A finite state space $\mathcal{S} = 2^{\{1, \dots, n\}}$ Set of all true features
2. A finite action space $\mathcal{A} = \Omega$ Actions = operators
3. An initiability action function $I(s, \omega) = \text{pre}(\omega) \subseteq s$ Preconditions hold
4. A transition function $F(s, \omega) = (s - \text{del}(\omega)) \cup \text{add}(\omega)$ Effects
5. A cost function $C(s, \omega, s') = 1$ For simplicity
6. An initial state $s_0 \in \mathcal{S}$
7. A goal function $G(s) = g \subseteq s$ g is another feature set

Lifted Operators

It is often convenient to define operators with **parameters**: placeholders for objects

Objects can also be **typed**

Preprocessing: *ground* all lifted operators with all combinations of objects (obeying types)

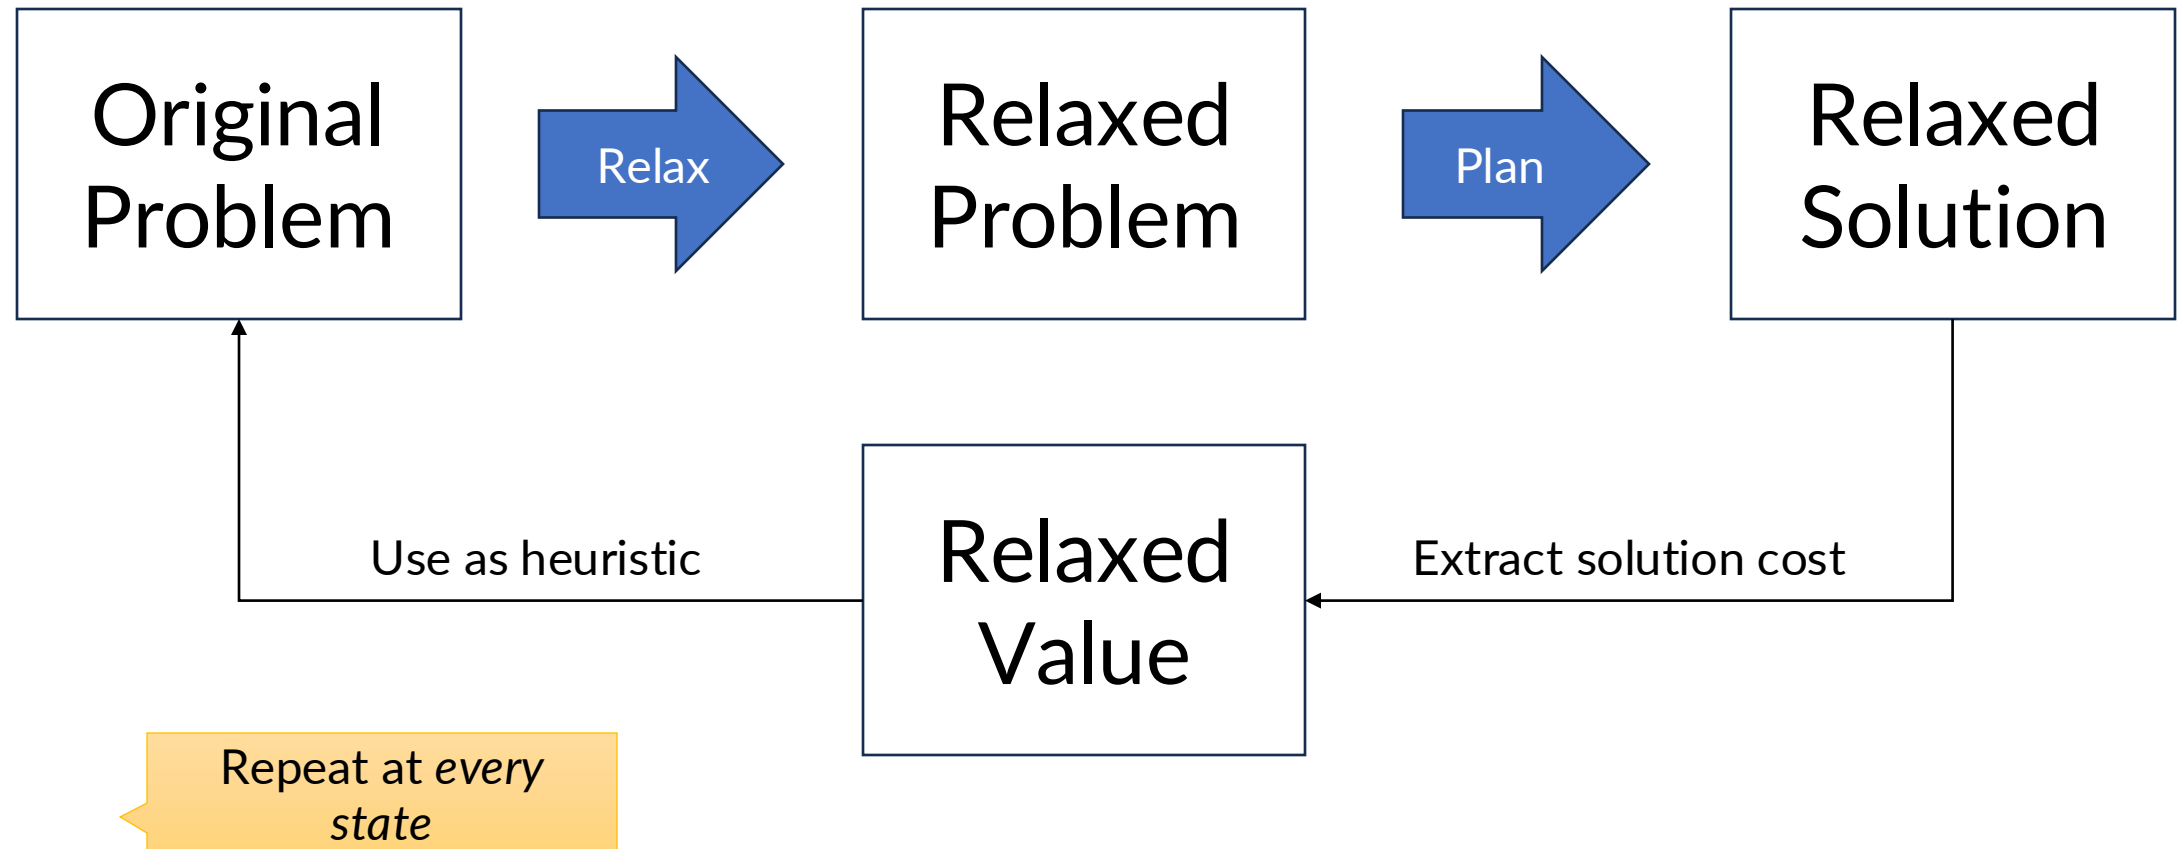
Pick(?x, ?y):

Preconditions: {HandEmpty(),
 On(?x, ?y),
 Clear(?x)}

Add effects: {Holding(?x), Clear(?y)}

Delete effects: {HandEmpty(),
 On(?x, ?y)
 Clear(?x)}

A Recipe for Heuristic Generation



Delete Relaxation

Pick(?x, ?y):

Preconditions: {HandEmpty(),
On(?x, ?y),
Clear(?x)}

Add effects: {Holding(?x),
Clear(?y)}

Delete effects: {HandEmpty(),
On(?x, ?y)
Clear(?x)}



Relax

Pick(?x, ?y):


Preconditions: {HandEmpty(),
On(?x, ?y),
Clear(?x)}

Add effects: {Holding(?x),
Clear(?y)}

Delete effects: {}

Delete-Relax: Our Second Problem-Derived Heuristic

The delete-relax heuristic $V_{DR}(s)$ is the optimal cost of the **relaxed** planning problem with initial state s .





Goal: holding(a) & holding(b)

What is $V_{DR}(s)$?
What is $V^*(s)$?

Is V_{DR} admissible?



Delete Relaxation Can Help!



But these plots are extremely misleading. Why?

More Revealing Plots

Solving delete-relaxed problems exactly is formally hard (NP-complete)



hFF: A Better Delete Relaxation Heuristic

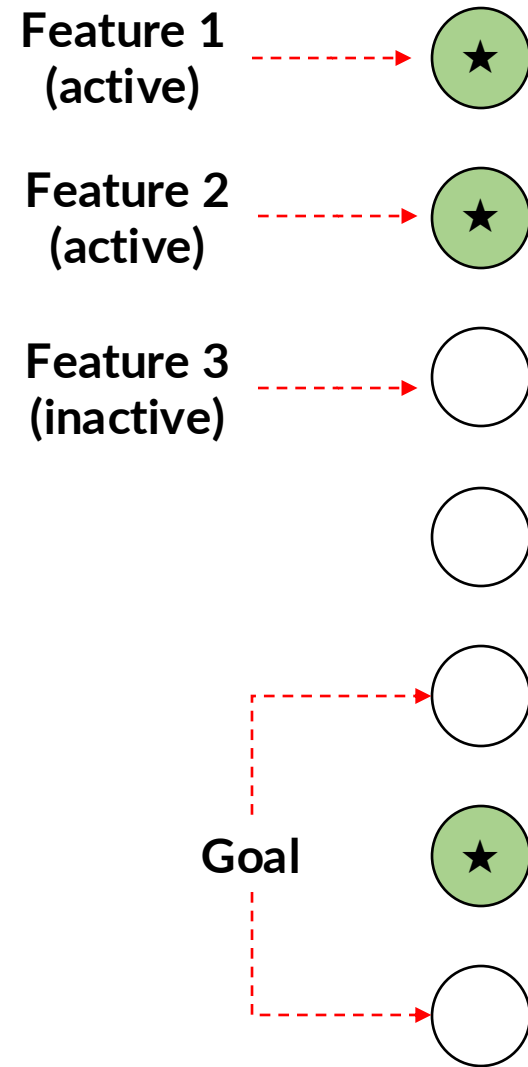
Construct a *non-optimal* relaxed plan in a particular way:





Forward pass:

1. Imagine we could execute *all* initiatable actions simultaneously
2. Aggregate the next states into superset of all active features
3. Repeat (1) and (2) until convergence (or goal is active)


Backward pass:

1. Build a relaxed plan by selecting “necessary” actions




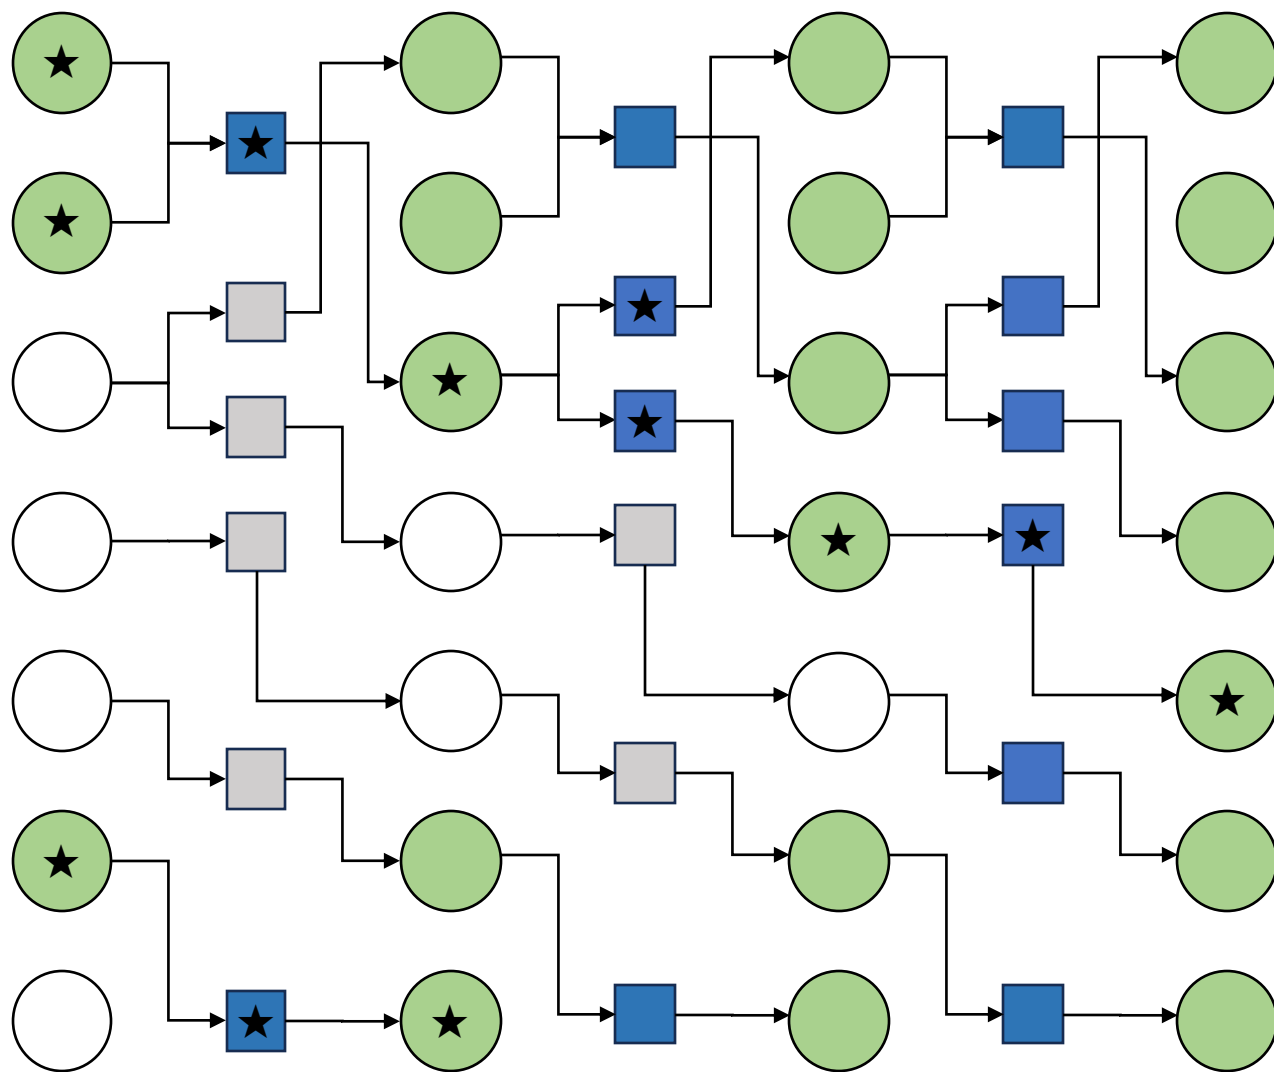


Operators

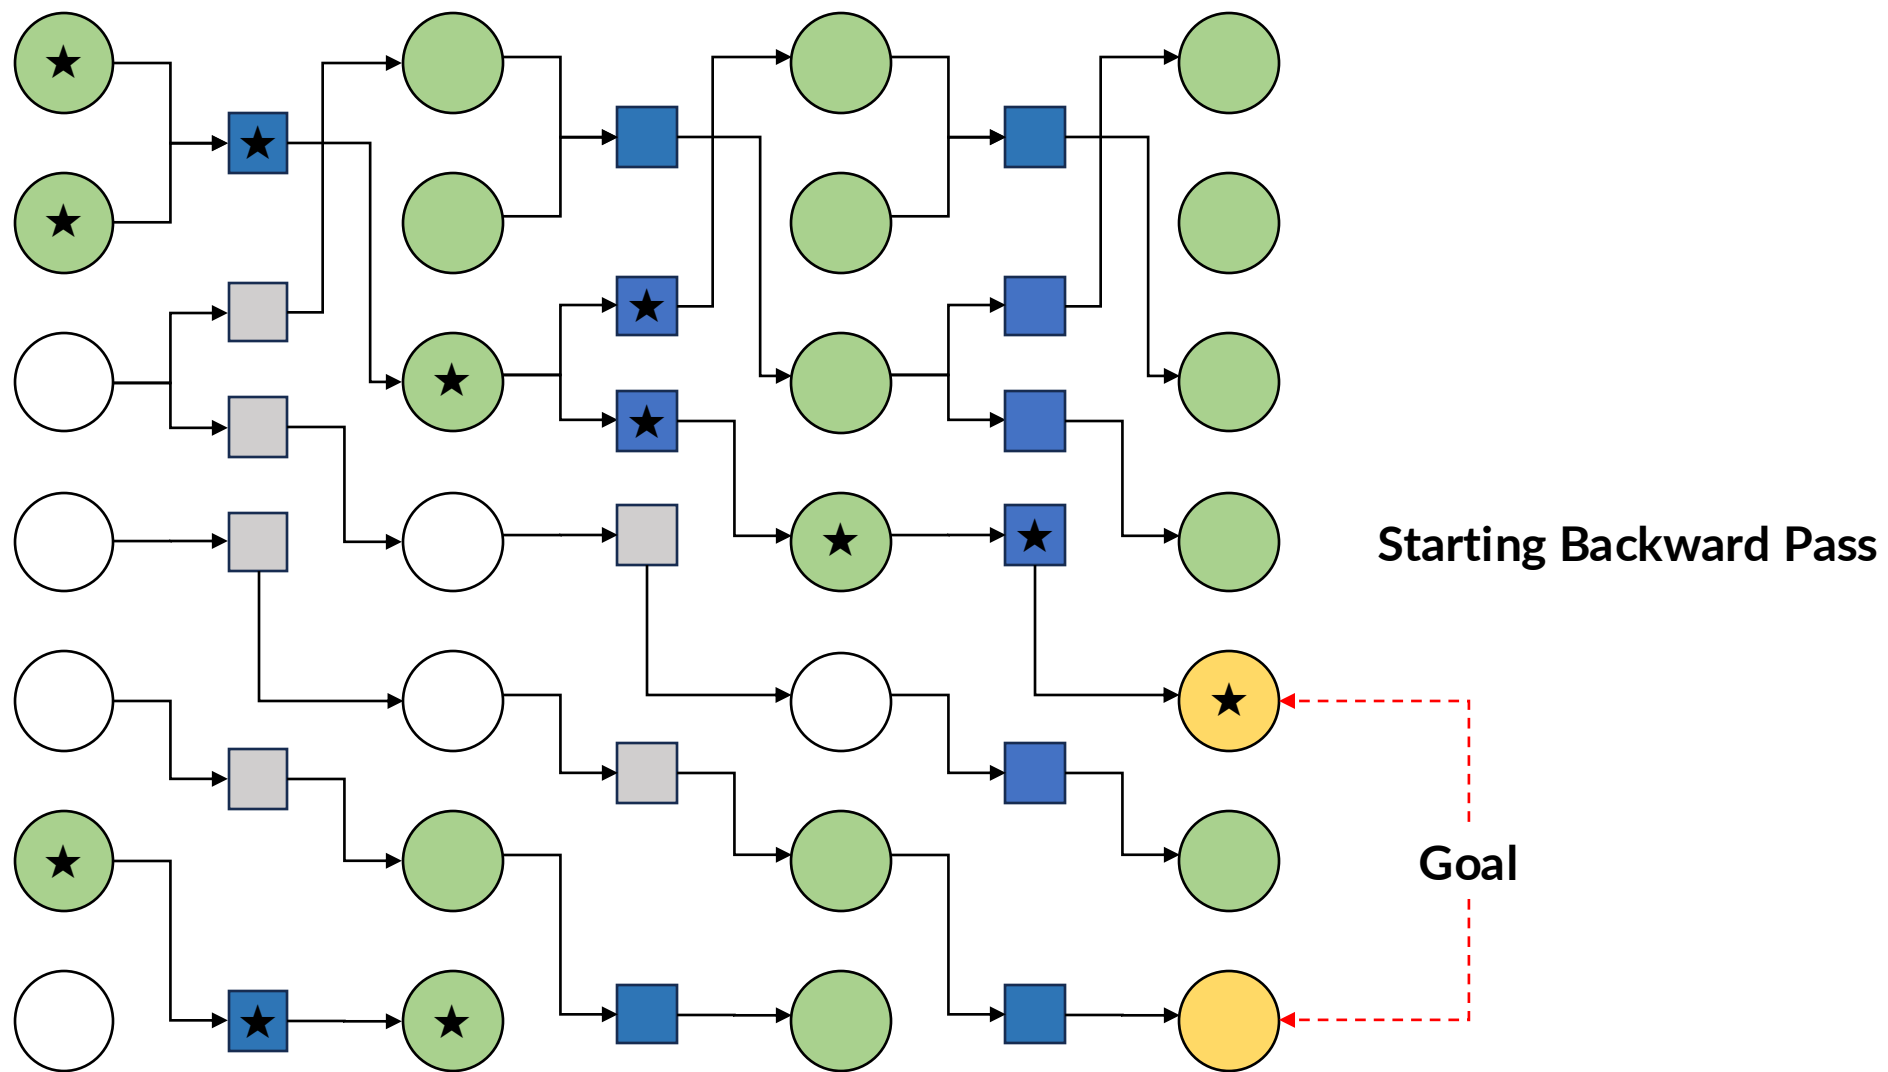


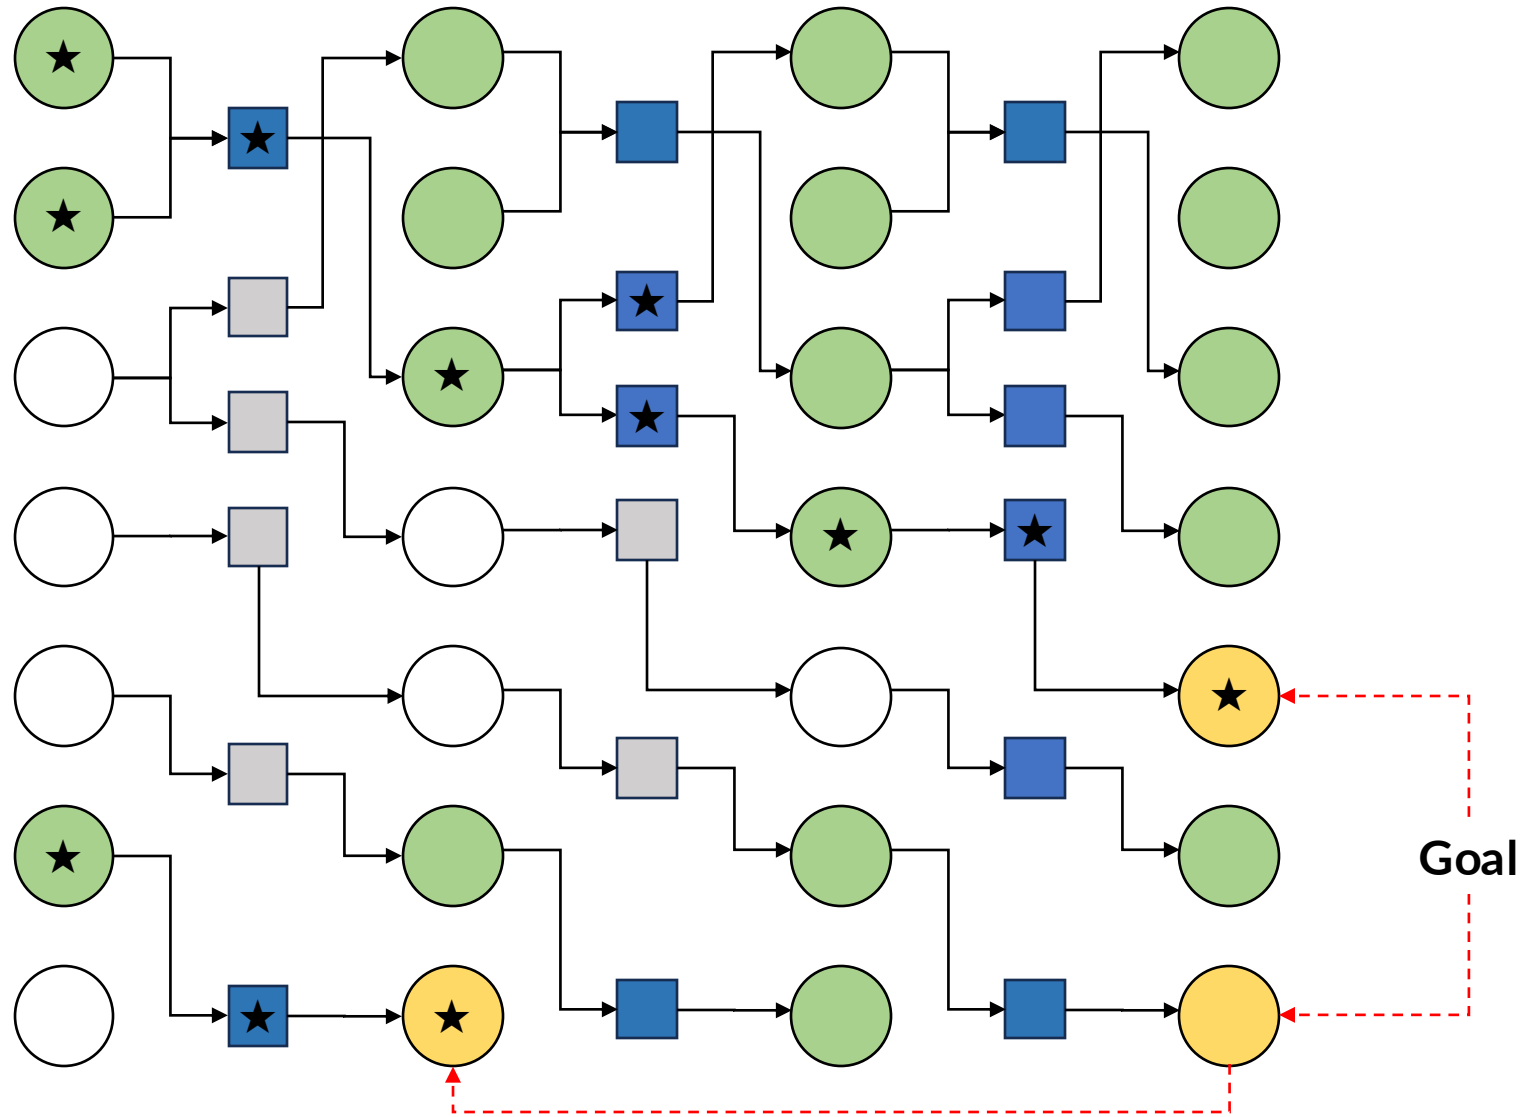
Operators

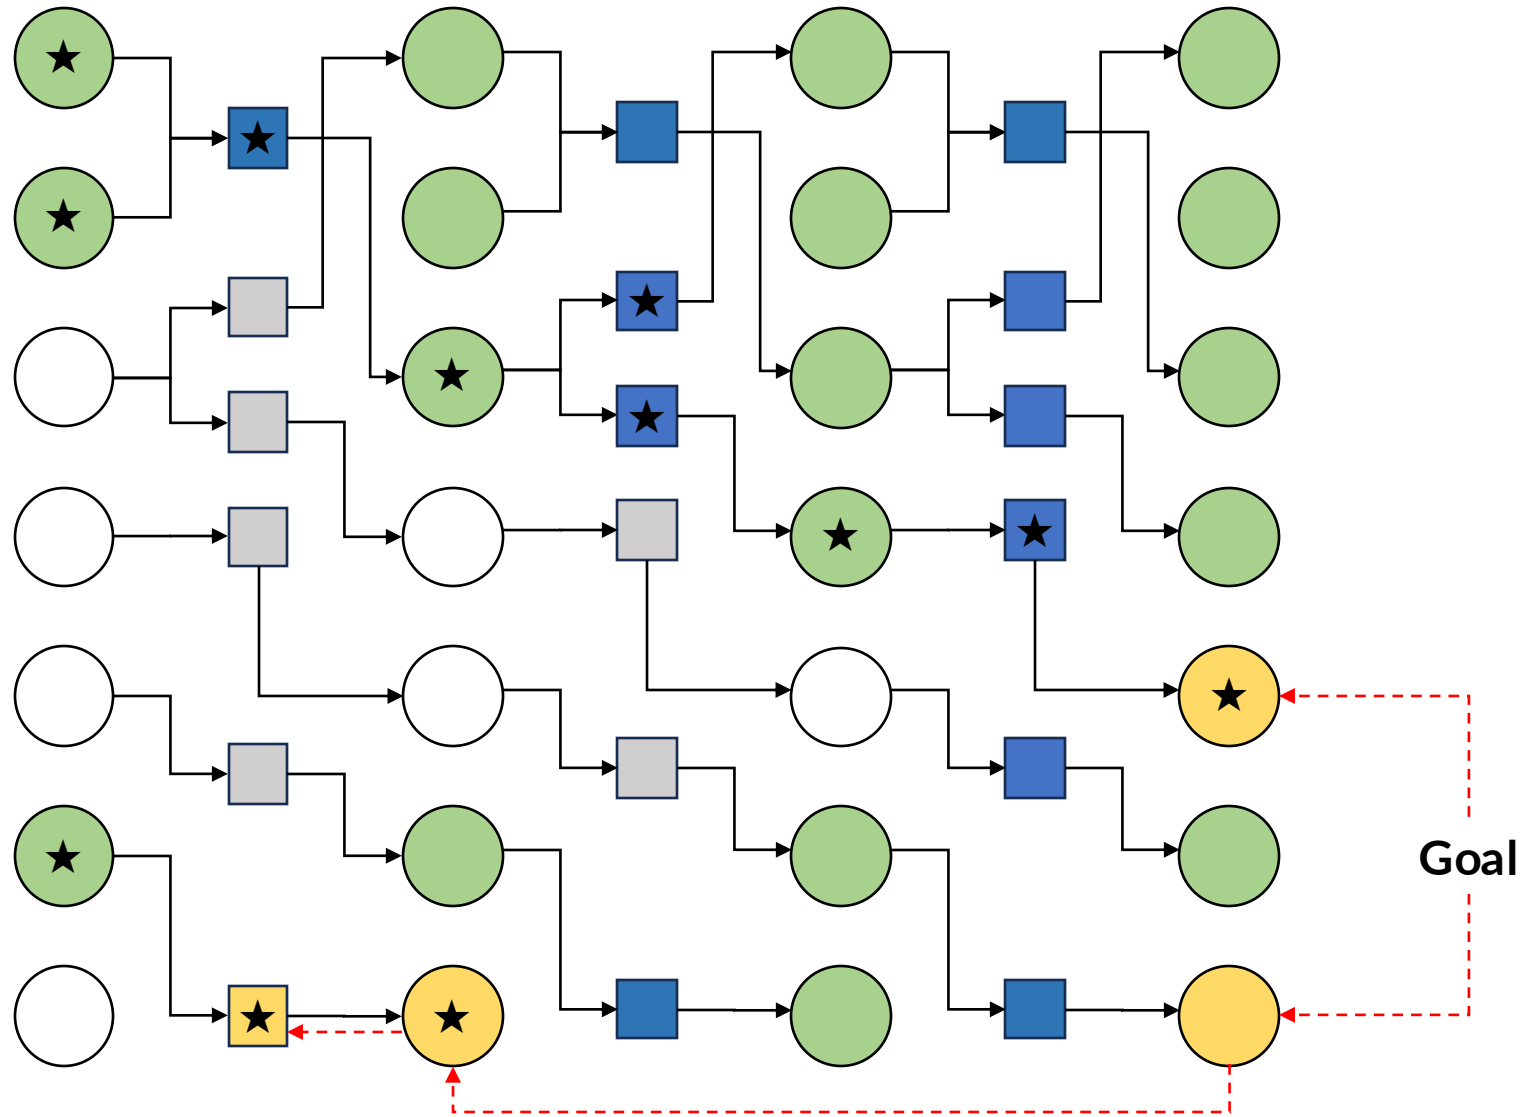


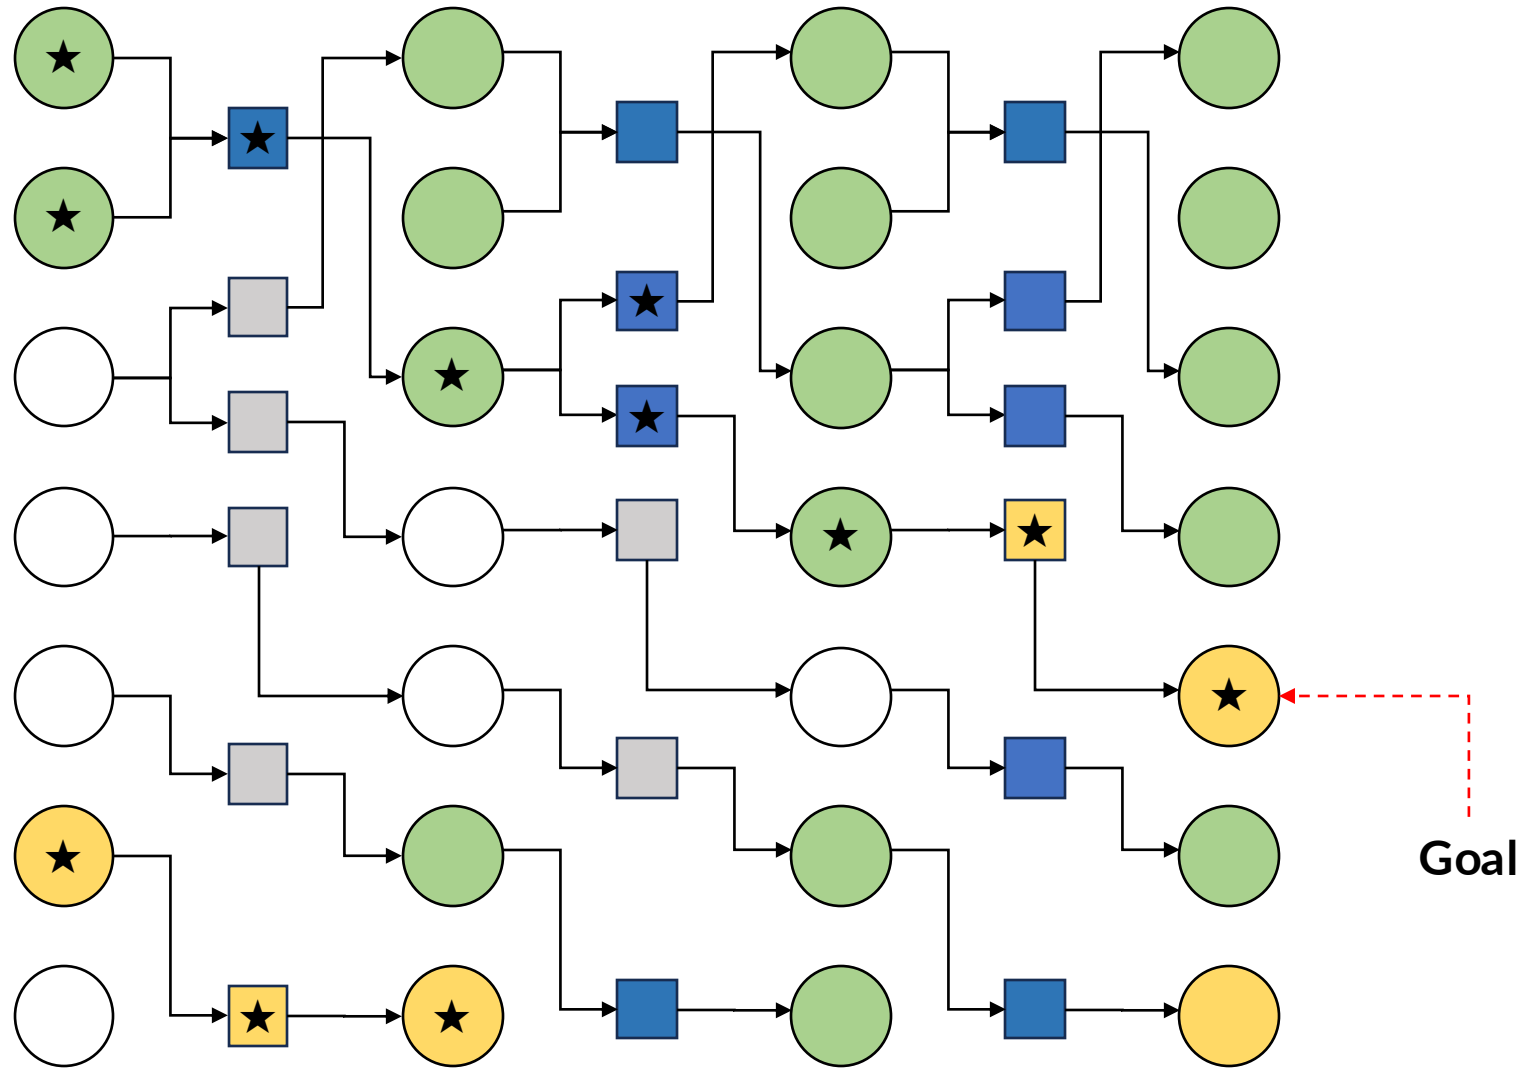


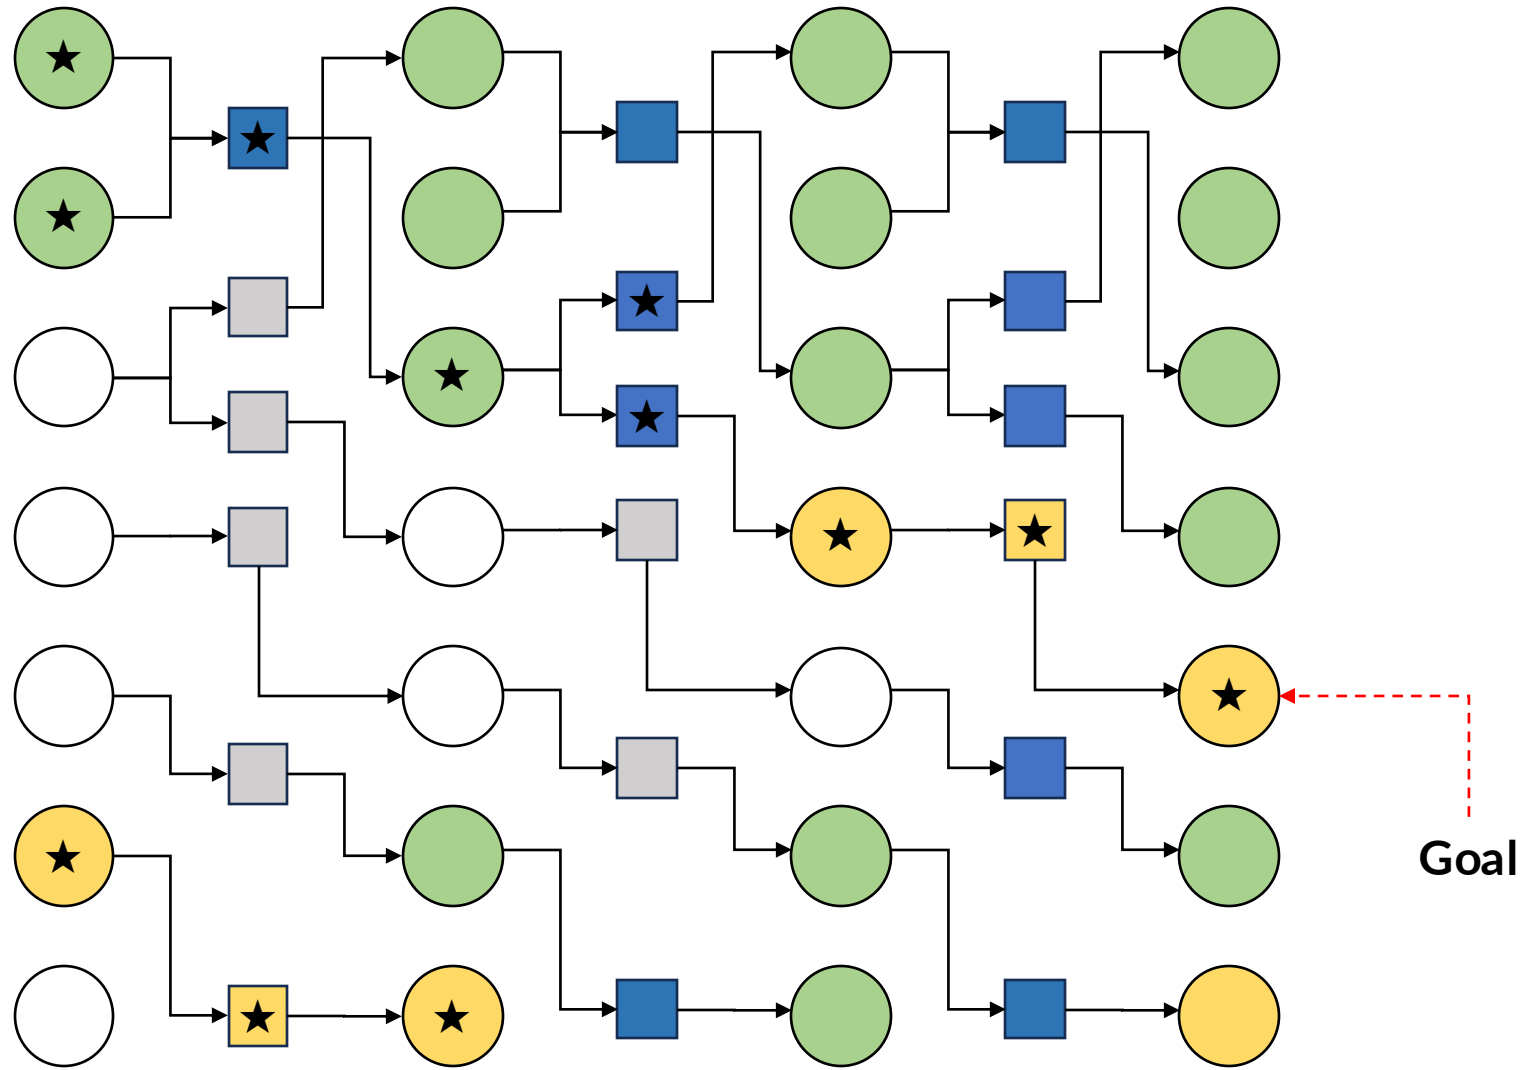
Convergence
Forward Pass Complete

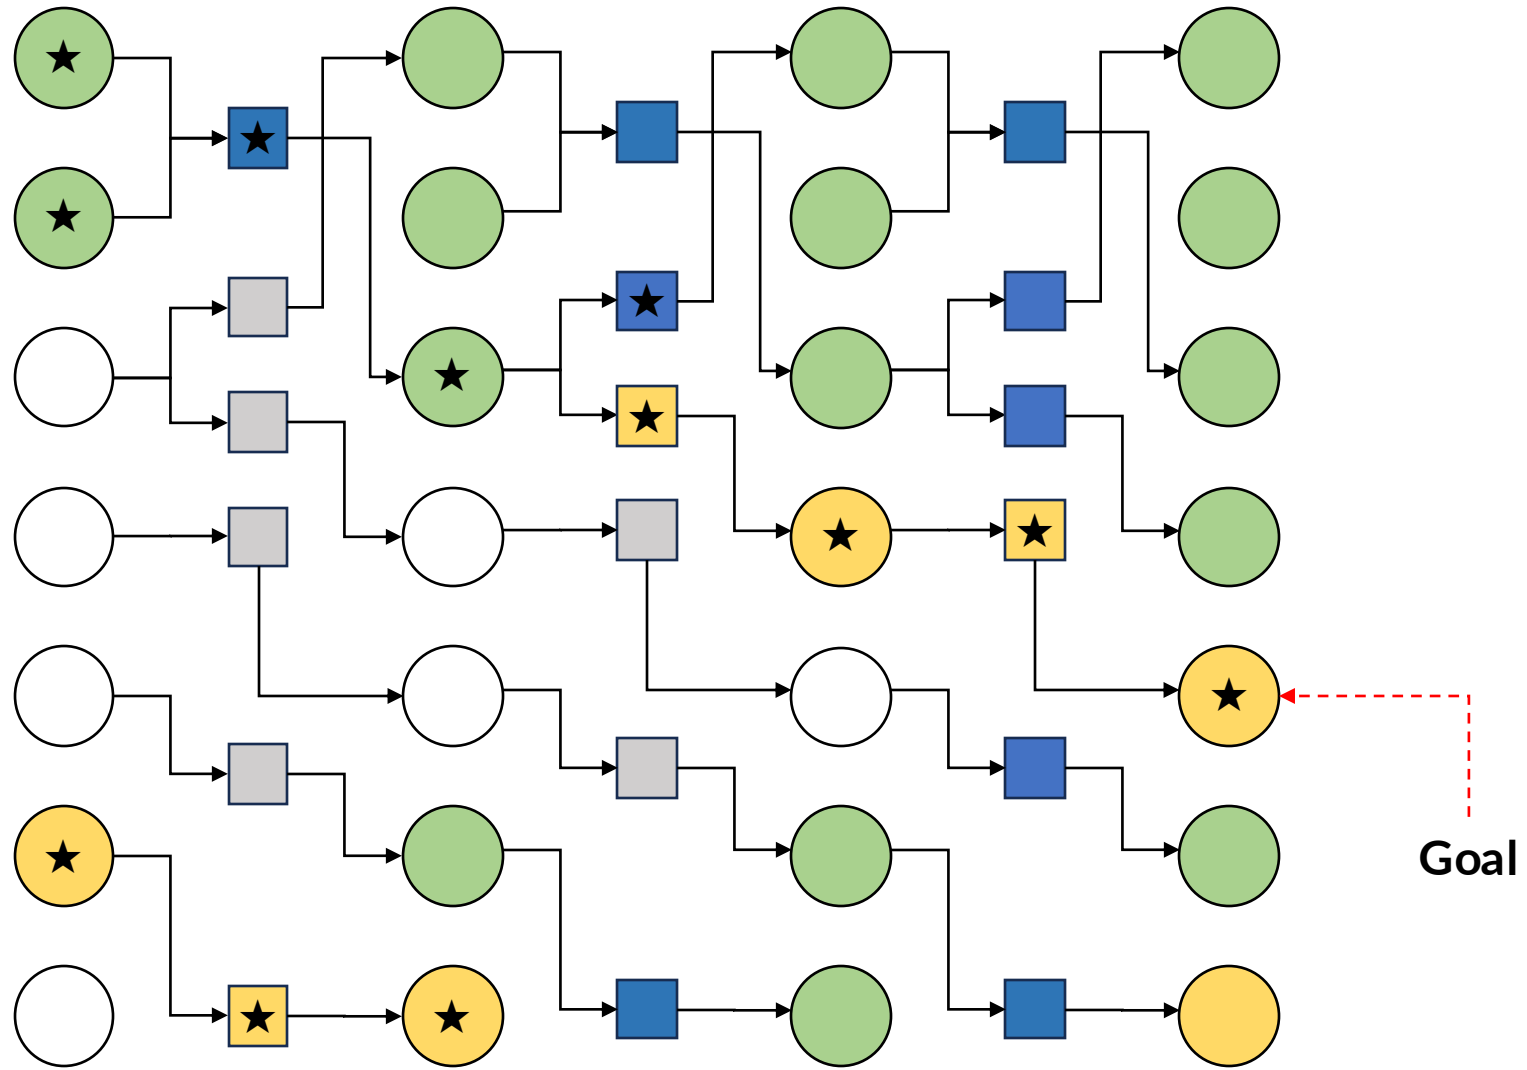


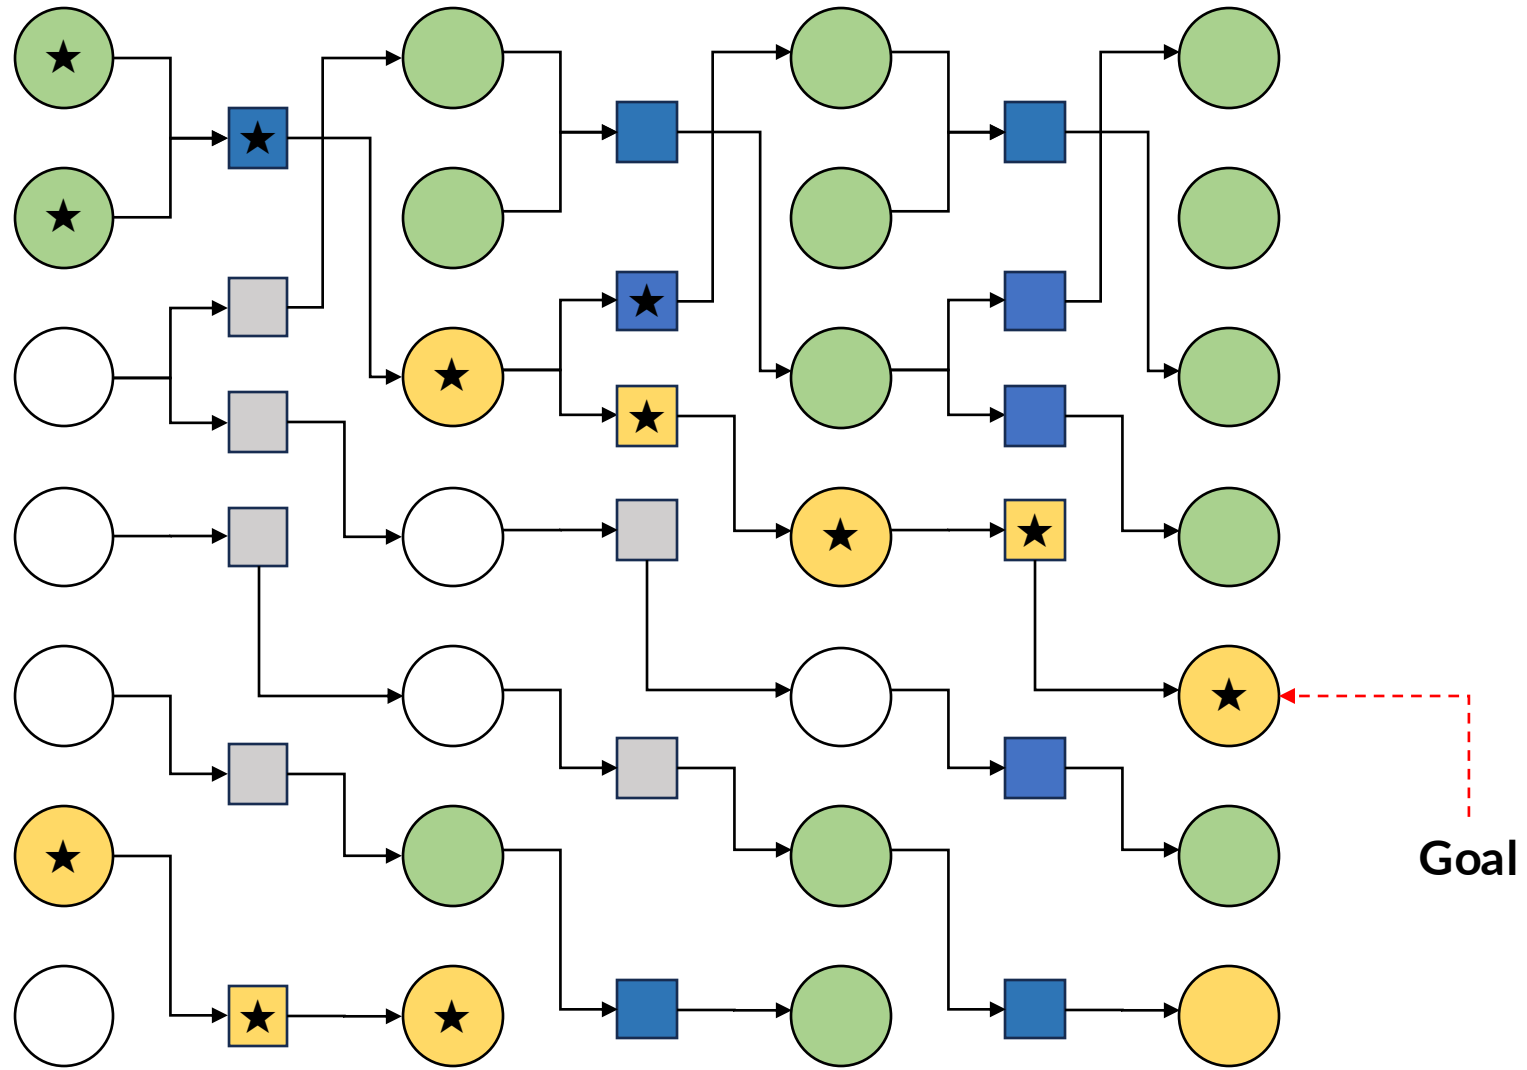


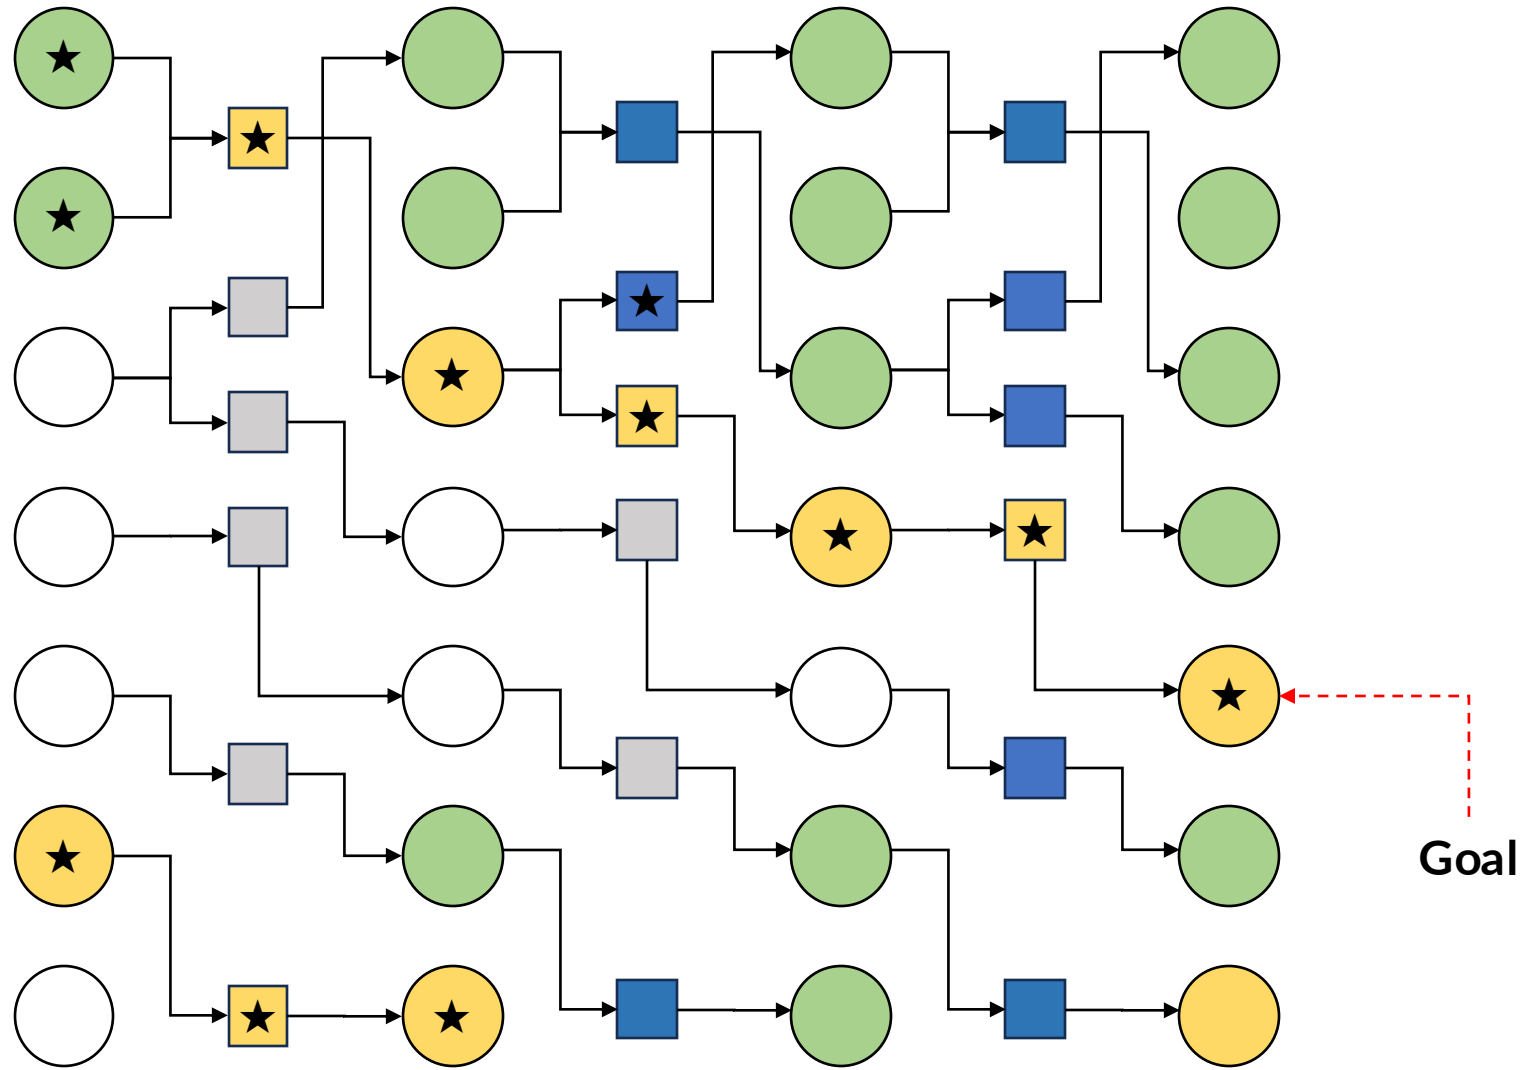











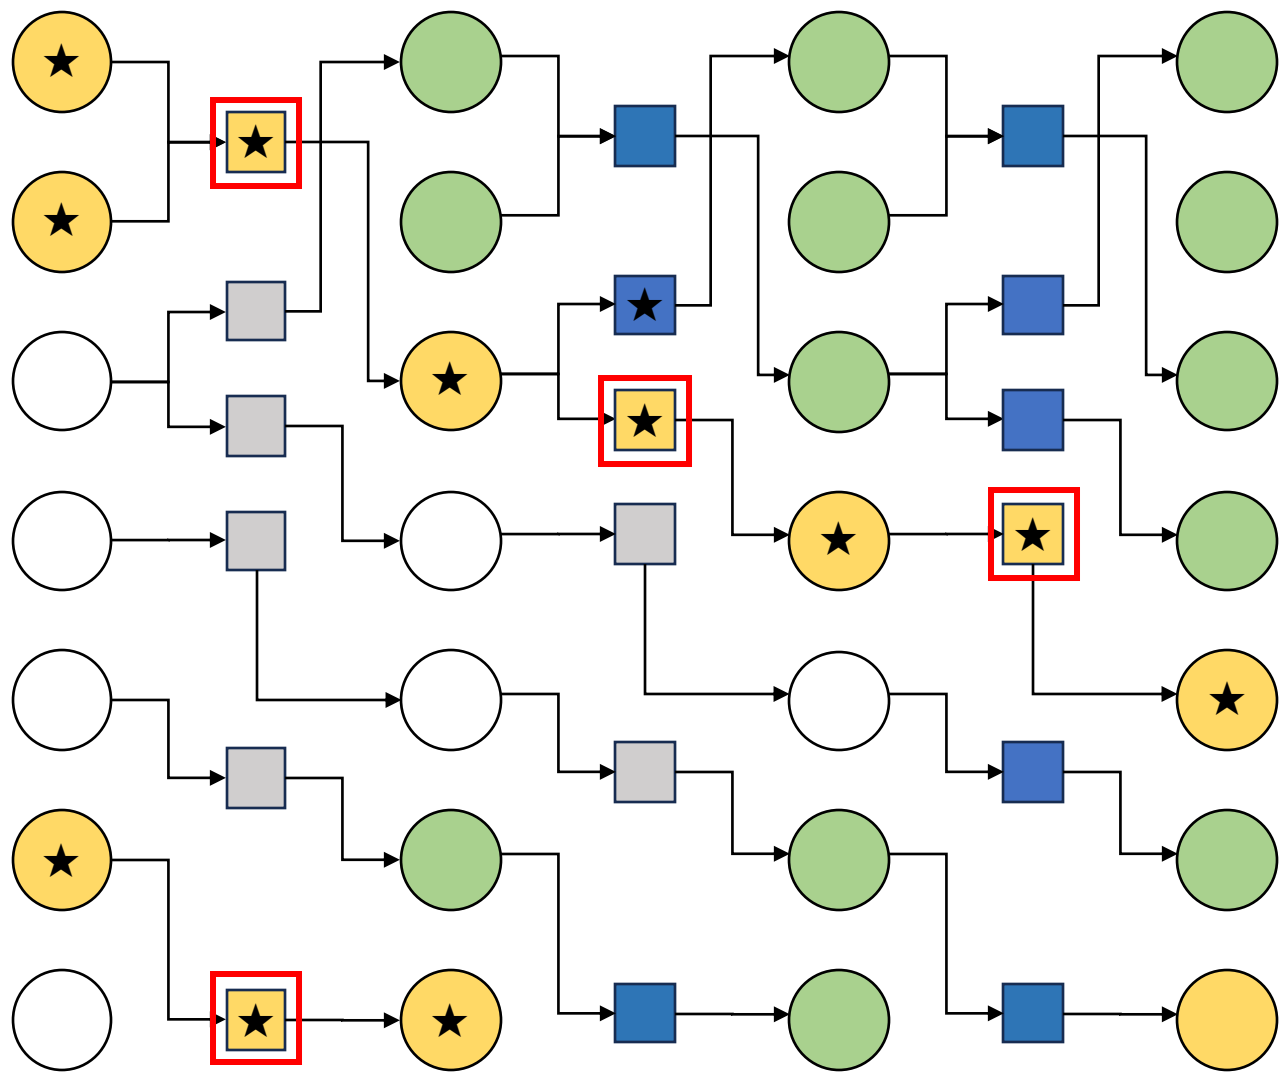


Relaxed
Plan



Relaxed
Plan

$$V_{hFF} = 4$$

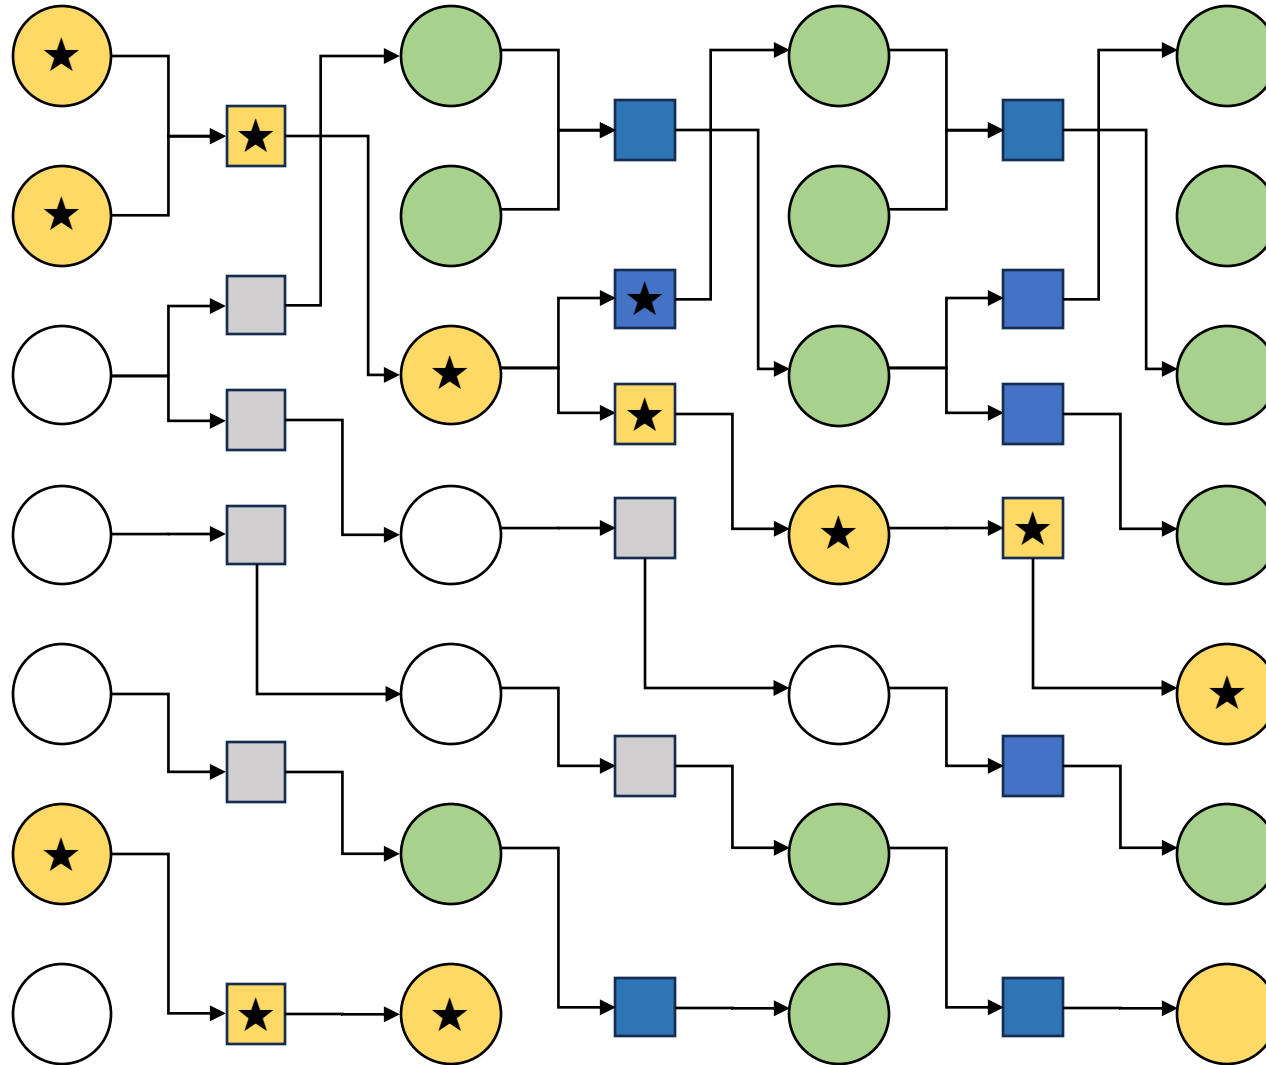


Terminology: Relaxed Planning Graph



Also useful for
testing *reachability*

Also useful for
pruning actions



hFF Can Really Help!

